

AI Front Desk for early education

Parent demo: brightwheel-front-desk.pages.dev | Operator: `/operator`

A working prototype of an out-of-the-box AI Front Desk. A parent asks a question by text or voice, in English or Spanish, and gets a fast, center-specific answer they can trust. The operator owns the source of truth, sees where the bot struggled, and improves it in one tap. Built for a fictional center (Cottonwood Sprouts, Albuquerque) with invented data.

The problem

Administrators lose hours a day answering the same questions: hours, tuition, illness policy, lunch, tours. The parent on the other end is anxious and busy, and a handbook PDF is the opposite of helpful at 7am with a feverish toddler. The job is to answer in a way a worried parent can trust, without adding work for a stretched operator.

Try it in about 90 seconds

- Tap the parent starters ("Veterans Day?", "infant tuition?", "what's for lunch today?"). Note the **source chip** and **confidence dot**.
- Ask in Spanish ("cuanto cuesta para bebes?"). It answers in Spanish.
- Ask "my child has a fever, can they come in?" It never gives medical advice; it shares policy and routes to a person.
- Ask something off-book ("do you have a pool?"). It says it doesn't have that on file instead of guessing.
- Open `/operator`: the **Impact** view, the **urgent-first inbox**, **Gap Radar** (teach the bot in one tap), and **Import from a photo**.

How it earns trust

- **Grounded and cited.** Answers only from the center handbook, with the source named; it says plainly when something is not on file.
- **A safety net that is code, not a prompt.** Medical, medication, emergency, and child-safety wording is caught deterministically before the model, in English and Spanish, and escalated to a person or 911.
- **Graceful uncertainty.** Low-confidence and sensitive replies escalate to a human and are logged as gaps.

What makes it more than a chatbot

- **It teaches itself.** Gap Radar drafts a handbook entry from each unanswered question; the operator approves in one tap. Health topics are held for manual review.
- **It onboards in minutes.** Snap a photo of a paper handbook and a vision model drafts editable sections.
- **It shows its value.** The operator opens to estimated hours saved and answer rate.

Under the hood

Cloudflare end to end (Pages, Functions, D1) with Groq (Llama) as the language model and Cloudflare Workers AI as an automatic fallback, so it is fast, cheap, and never fully dark. The whole handbook fits in the prompt, so there is no vector database. The model provider lives behind one function. The build was peer-reviewed at each stage with a second AI (Codex).

Notes

Fictional data; no personal data used. The live demo runs on a free, disposable Groq key that expires about 30 days from setup; when it lapses, answers fall back to Workers AI automatically and only photo import pauses.